

Precise Exception: Keeping sequential execution with pipeline

Seehwan Yoo

Dankook University

Dept. of Mobile Systems Engineering

Review: Pipelined u-Arch

- Divide instruction processing into stages
 - Fetch, Decode, Execute, Memory, Writeback
 - Could be much more
- Each stage utilizes different hardware (logic circuit)
 - IMem, Decoder/RegFile, ALU/adder, DMem, RegFile
- Why not utilize them all at the same time?
 - Overlap the execution of different stages
 - Fill in different instructions into pipeline stages
 - Instruction-Level Parallelism
 - Parallel execution of multiple instructions

Pipeline in a real world

- Take care dependencies
- Data dependency
 - Some instructions have dependency
- Control dependency
 - Some instruction changes control flow
- Handling dependency
 - Stall (hardware-based interlocking)
 - Software-based interlocking (add NOP instruction)
 - Forwarding, Scoreboarding
 - Delayed branch
 - Branch prediction, Predicated execution
 - Fine-grained multi-threading

Consider Multicycle Functional units!


- Integer add is much easier than floating-point add
 - And multiply, and memory operations
- Integer execution 1 cycle
- FPMultiply execution 8 cycles
- Say, Memory 10 cycles
- Consume cycles until the result is produced

FMUL R4 \leftarrow R1, R2

ADD R3 \leftarrow R1, R2

FMUL R2 \leftarrow R5, R6

ADD R4 \leftarrow R5, R6




Sequential execution is broken!

- Completes instructions out-of-order!
 - Note that fetch instructions in-order
- How would you debug?
- What if exception/interrupt occurs?
- Need to preserve ordered completion
- Commit == graduate == retire

A non-fix

- Simple dumb idea



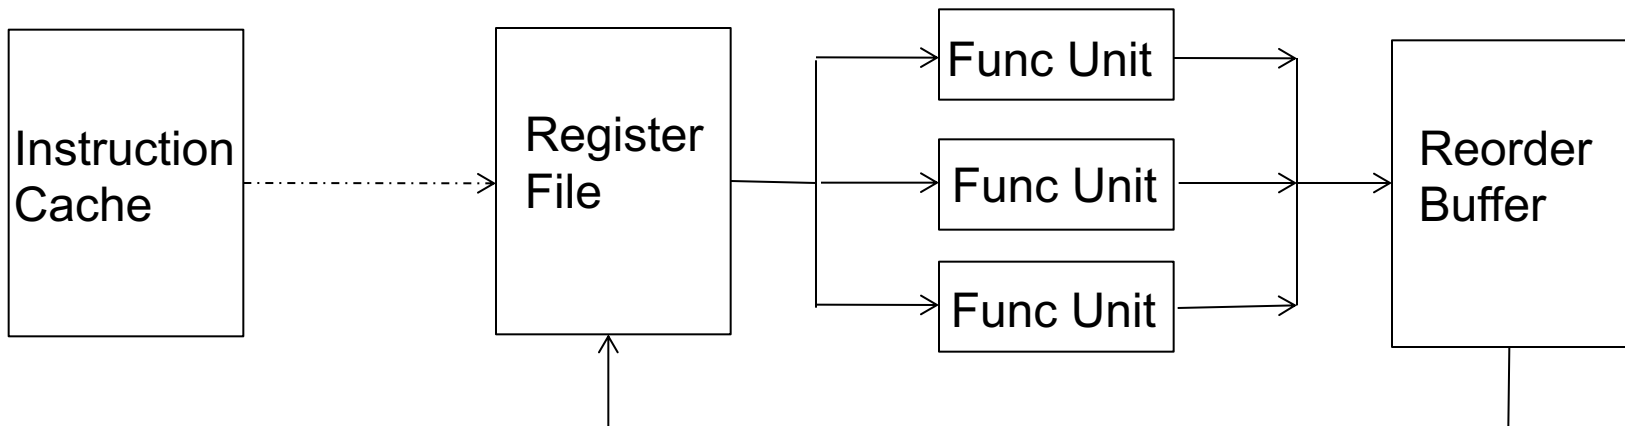
- What about memory operation?
 - Each functional unit takes 500 cycles?

Some solutions

- Reorder buffer
- History buffer
- Future register file
- Checkpointing
- Recommended Reading
 - Smith and Plezskun, “[Implementing Precise Interrupts in Pipelined Processors](#),” IEEE Trans on Computers 1988 and ISCA 1985.
 - Hwu and Patt, “[Checkpoint Repair for Out-of-order Execution Machines](#),” ISCA 1987.

ROB (Reorder-buffer)

- Idea: Complete instructions out-of-order, but reorder them before making results visible to architectural state
- When instruction is decoded it reserves an entry in the ROB
- When instruction completes, it writes result into ROB entry
- When instruction oldest in ROB and it has completed without exceptions, its result moved to reg. file or memory



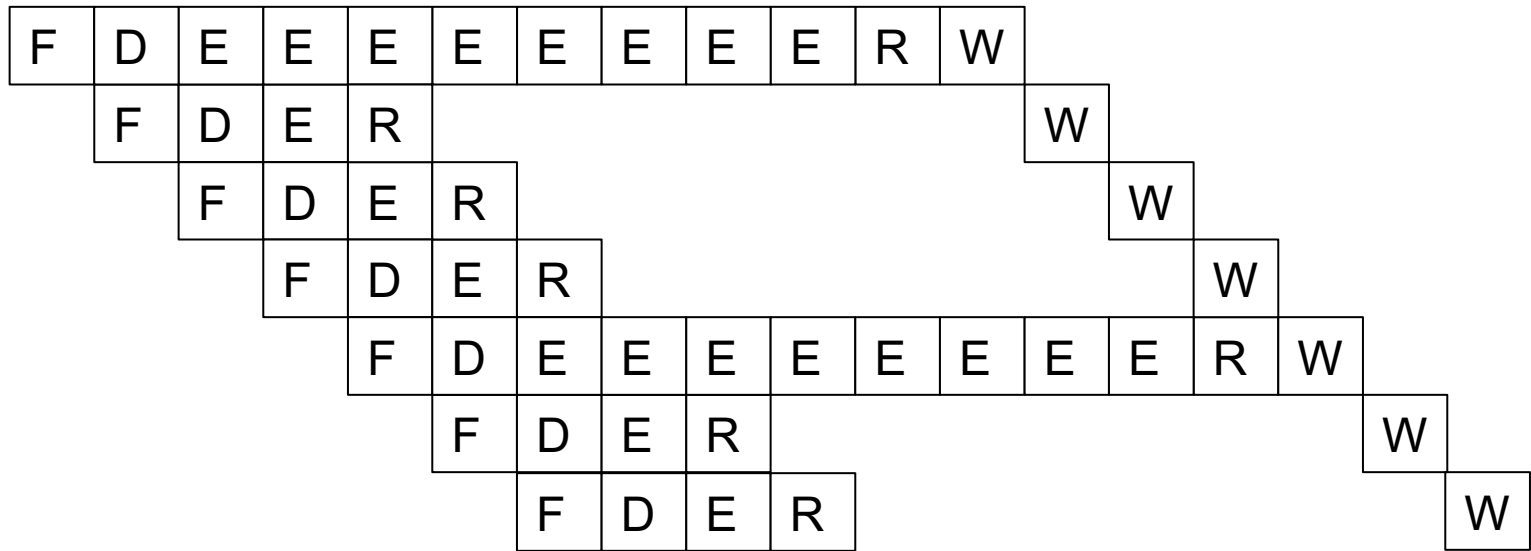
ROB entry

V	DestRegID	DestRegVal	StoreAddr	StoreData	PC	Valid bits for reg/data + control bits	Exc?
---	-----------	------------	-----------	-----------	----	---	------

- Need valid bits to keep track of readiness of the result(s)

ROB: independent ops


- Results first written to ROB, then to register file at commit time



- What if a later operation needs a value in the reorder buffer?
 - Read reorder buffer in parallel with the register file. **How?**

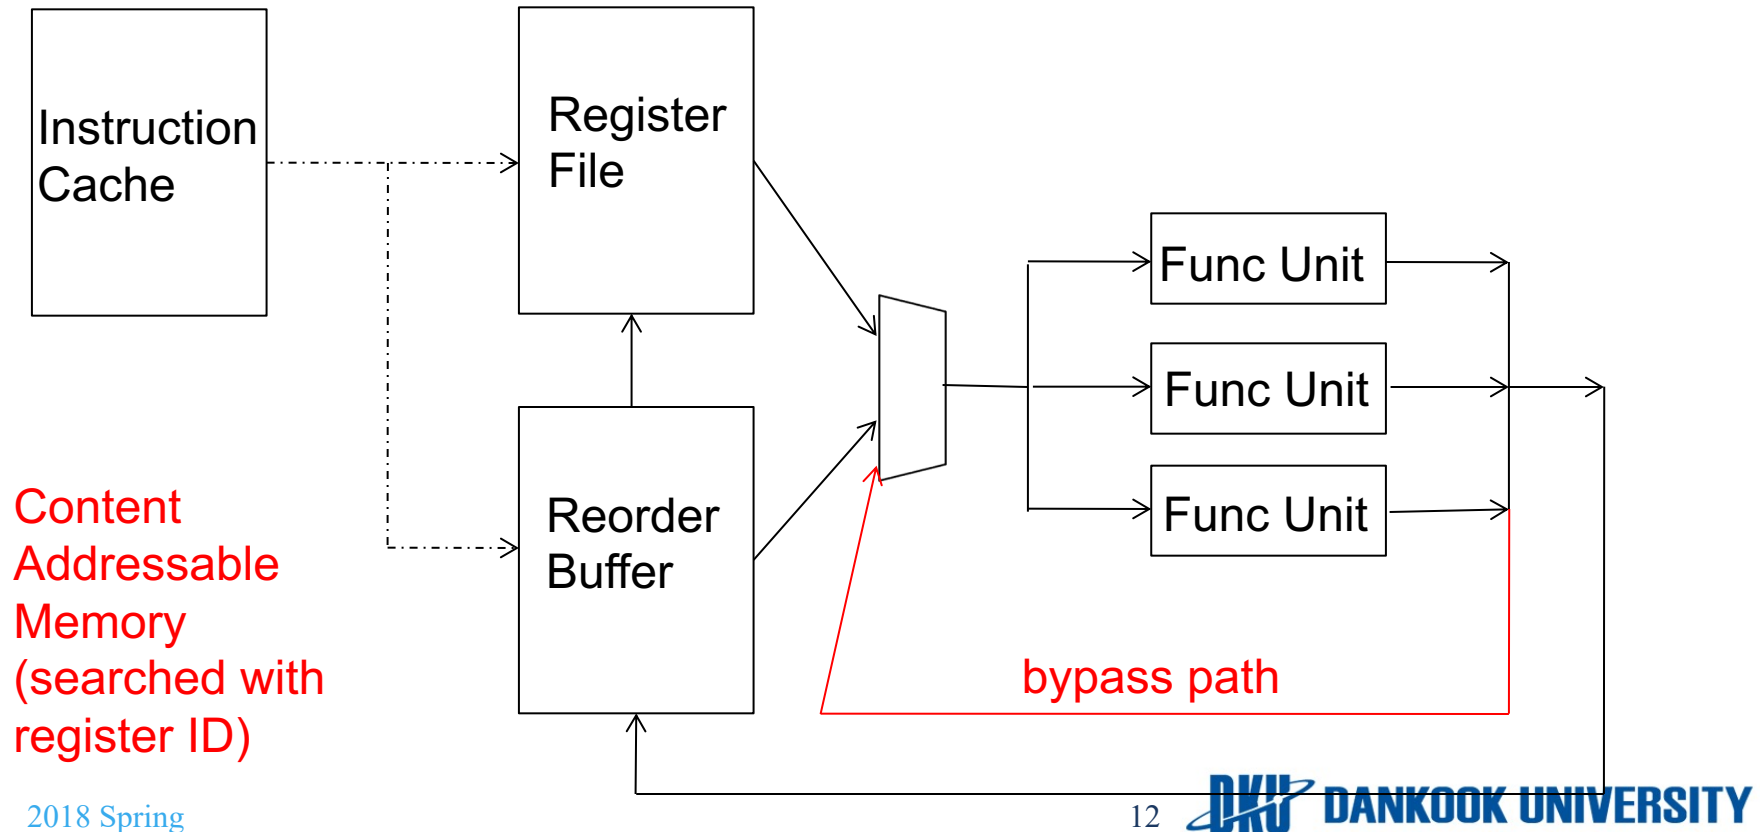
In-order pipeline with ROB

- **Decode (D)**: Access regfile/ROB, allocate entry in ROB, check if instruction can execute, if so **dispatch** instruction
- **Execute (E)**: Instructions can complete out-of-order
- **Completion (R)**: Write result to reorder buffer
- **Retirement/Commit (W)**: Check for exceptions; if none, write result to architectural register file or memory; else, flush pipeline and start from exception handler
- **In-order dispatch/execution, out-of-order completion, in-order retirement**



What if \exists data dependency?

- Now, register file has no recent data!
- Reorder buffer has updated value!
- Or you can get value from forwarding path (bypass network)



How to find value in ROB?

- Idea: Use indirection
- Access register file first
 - If register not valid, register file stores the ID of the reorder buffer entry that contains (or will contain) the value of the register
 - Mapping of the register to a ROB entry
- Access reorder buffer next
- What is in a reorder buffer entry?

V	DestRegID	DestRegVal	StoreAddr	StoreData	PC/IP	Control/valid bits	Exc?
---	-----------	------------	-----------	-----------	-------	--------------------	------

- Can it be simplified further?

Register renaming

- Output and anti dependencies are not true dependencies
 - **WHY? The same register refers to values that have nothing to do with each other**
 - **They exist due to lack of register ID's (i.e. names) in the ISA**
- The register ID is **renamed** to the reorder buffer entry that will hold the register's value
 - Register ID → ROB entry ID
 - Architectural register ID → Physical register ID
 - After renaming, ROB entry ID used to refer to the register
- This eliminates anti- and output- dependencies
 - Gives the illusion that there are a large number of registers

ROB: +/-

- Advantages

- Conceptually simple for supporting precise exceptions
- Can eliminate false dependencies

- Disadvantages

- Reorder buffer needs to be accessed to get the results that are yet to be written to the register file
 - CAM or indirection → increased latency and complexity

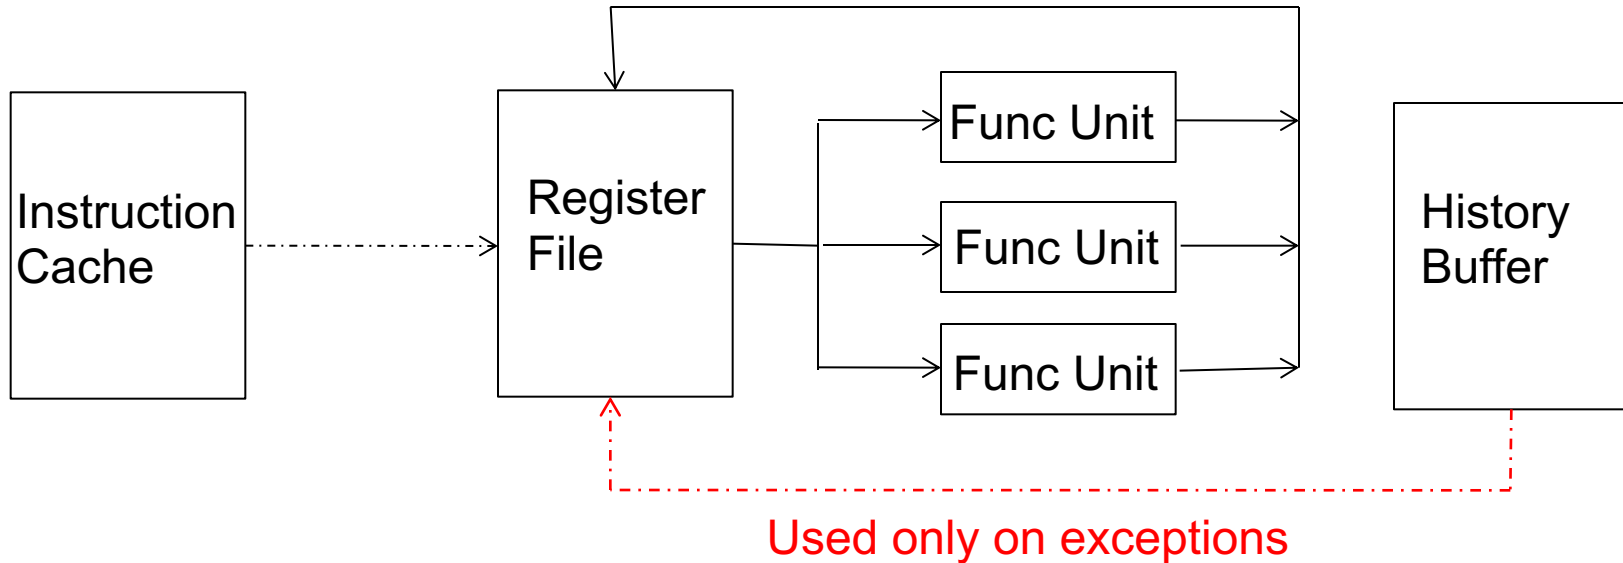
- Other solutions aim to eliminate the cons

- History buffer
- Future file
- Checkpointing

History Buffer

- Idea: Update the register file when instruction completes, but UNDO UPDATES when an exception occurs
- When instruction is decoded, it reserves an HB entry
- When the instruction completes, it stores the old value of its destination in the HB
- When instruction is oldest and no exceptions/interrupts, the HB entry discarded
- When instruction is oldest and an exception needs to be handled, old values in the HB are written back into the architectural state from tail to head

History Buffer

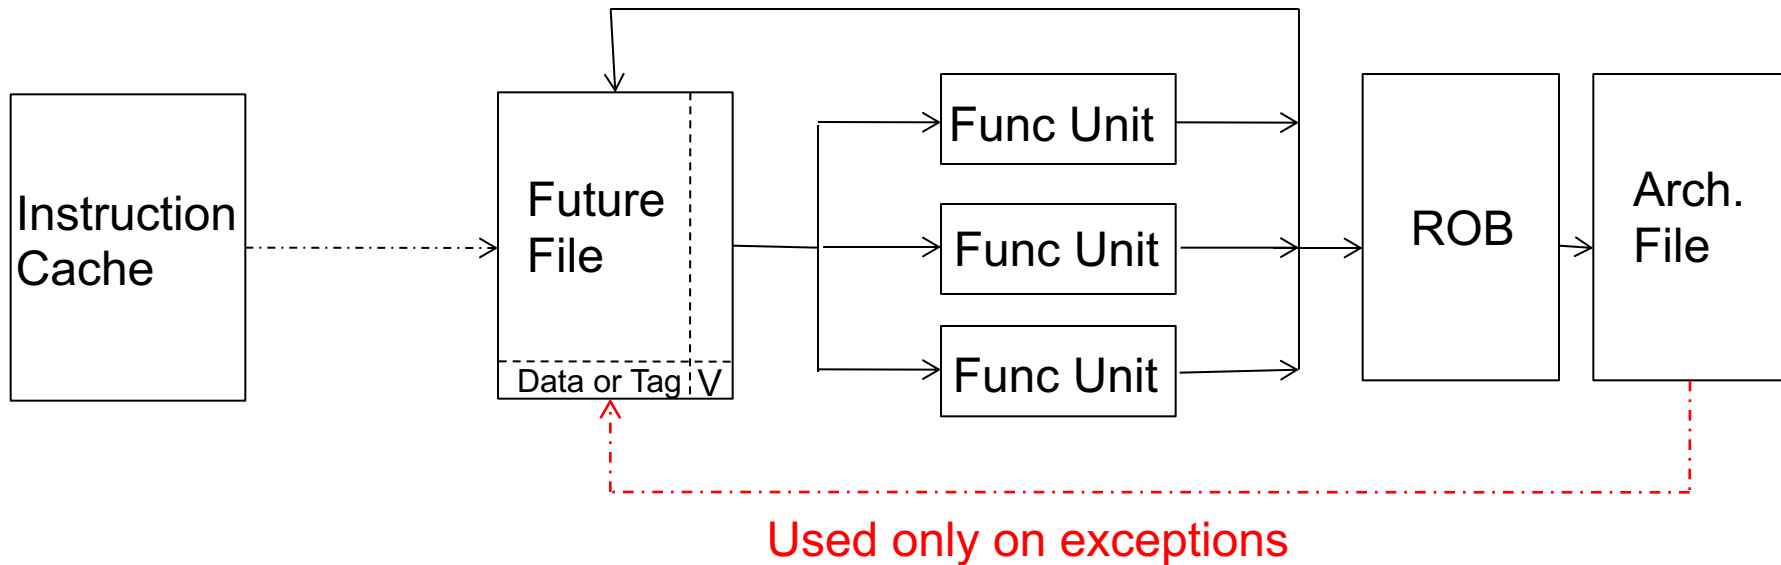


- **Advantage:**
 - Register file contains up-to-date values. History buffer access not on critical path
- **Disadvantage:**
 - Need to read the old value of the destination register
 - Need to unwind the history buffer upon an exception → increased exception/interrupt handling latency

Future File + ROB

- Idea: **Keep two register files (speculative and architectural)**
 - Arch reg file: Updated in program order for precise exceptions
 - Use a reorder buffer to ensure in-order updated
 - Future reg file: Updated as soon as an instruction completes (if the instruction is the youngest one to write to a register)
- Future file is used for fast access to latest register values (speculative state)
 - Frontend register file
- Architectural file is used for state recovery on exceptions (architectural state)
 - Backend register file

Future File



■ Advantage

- ❑ No need to read the values from the ROB (no CAM or indirection)

■ Disadvantage

- ❑ Multiple register files
- ❑ Need to copy arch. reg. file to future file on an exception

Side Note: Branch misprediction

- Similar to Exception
 - Except it is not visible to SW
- Much more common than exception, so state recovery has to be fast
- Reorder Buffer
 - Wait until branch is the oldest instruction in the machine
 - Flush pipeline afterwards
- History buffer
 - Undo all instructions after the branch by rewinding from the tail of the history buffer until the branch & restoring old values into the register file
- Future file
 - Wait until branch is the oldest instruction in the machine
 - Copy arch. reg. file to future file and flush pipeline

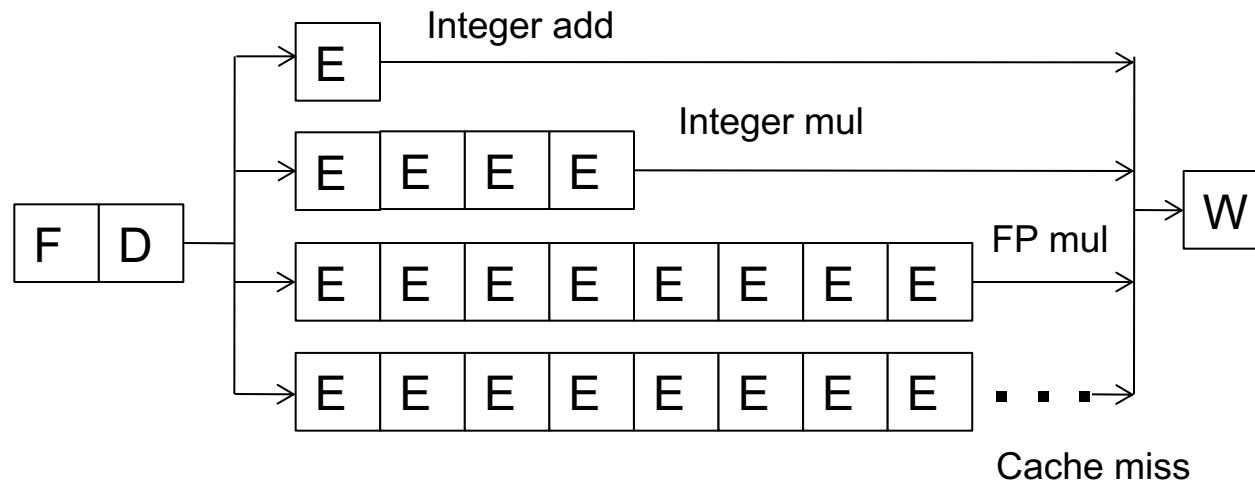
Side Note2: Store Buffer

- Handling out-of-order completion of memory operations
 - UNDOing a memory write more difficult than UNDOing a register write. **Why?**
 - **One idea:** Keep store address/data in reorder buffer
 - How does a load instruction find its data?
 - **Store/write buffer:** Similar to reorder buffer, but used only for store instructions
 - Program-order list of un-committed store operations
 - When store is decoded: Allocate a store buffer entry
 - When store address and data become available: Record in store buffer entry
 - When the store is the oldest instruction in the pipeline: Update the memory address (i.e. cache) with store data

Out-of-order Execution

Making irrationally intelligent HW

An In-order Pipeline



- Problem: A true data dependency stalls dispatch of younger instructions into functional (execution) units
- Dispatch: Act of sending an instruction to a functional unit

A motivational case for OoO

- What do the following two pieces of code have in common (with respect to execution in the previous design)?

IMUL R3 \leftarrow R1, R2

ADD R3 \leftarrow R3, R1

ADD R1 \leftarrow R6, R7

IMUL R5 \leftarrow R6, R8

ADD R7 \leftarrow R3, R5

LD R3 \leftarrow R1 (0)

ADD R3 \leftarrow R3, R1

ADD R1 \leftarrow R6, R7

IMUL R5 \leftarrow R6, R8

ADD R7 \leftarrow R3, R5


- Answer: First ADD stalls the whole pipeline!
 - ADD cannot dispatch because its source registers unavailable
 - Later **independent** instructions cannot get executed
- How are the above code portions different?
 - Answer: Load latency is variable (unknown until runtime)
 - What does this affect? Think compiler vs. microarchitecture

Out-of-order Execution (Dynamic Scheduling)

- Idea: Move the dependent instructions out of the way of independent ones
 - Rest areas for dependent instructions: Reservation stations
- Monitor the source “values” of each instruction in the resting area
- When all source “values” of an instruction are available, “fire” (i.e. dispatch) the instruction
 - Instructions dispatched in dataflow (not control-flow) order
- Benefit:
 - **Latency tolerance**: Allows independent instructions to execute and complete in the presence of a long latency operation

In-order vs. Out-of-order Dispatch

- In order dispatch:



IMUL R3 \leftarrow R1, R2


ADD R3 \leftarrow R3, R1

ADD R1 \leftarrow R6, R7

IMUL R5 \leftarrow R6, R8

ADD R7 \leftarrow R3, R5

- Tomasulo + precise exceptions:



- 16 vs. 12 cycles

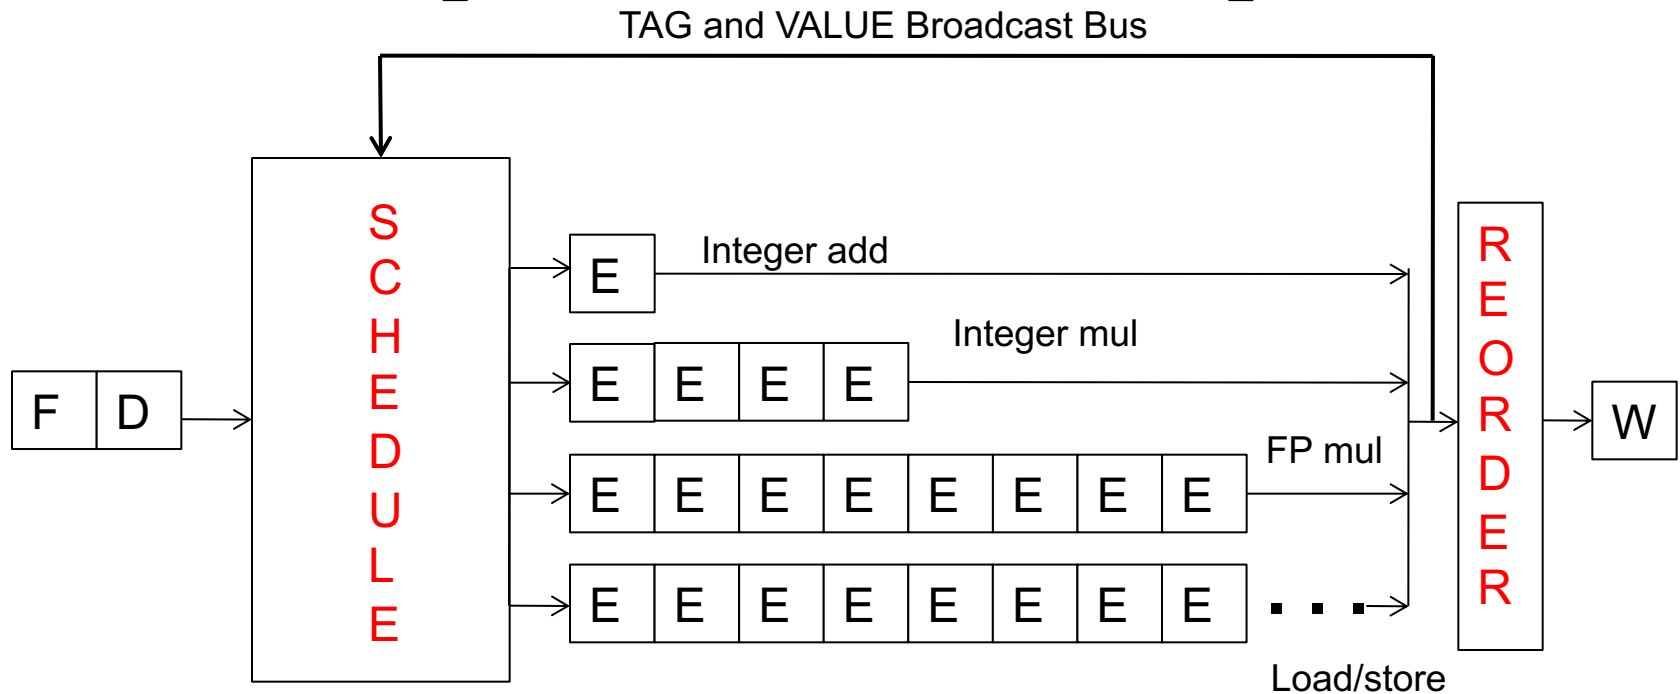
Enabling OoO Execution

1. Need to link the consumer of a value to the producer
 - **Register renaming:** Associate a “tag” with each data value
2. Need to buffer instructions until they are ready to execute
 - Insert instruction into **reservation stations** after renaming
3. Instructions need to keep track of readiness of source values
 - **Broadcast the “tag”** when the value is produced
 - Instructions **compare their “source tags”** to the broadcast tag → if match, source value becomes ready
4. When all source values of an instruction are ready, need to dispatch the instruction to its functional unit (FU)
 - What if more instructions become ready than available FUs?

Tomasulo's Algorithm

- OoO with register renaming invented by Robert Tomasulo
 - Used in IBM 360/91 Floating Point Units
 - **Read:** Tomasulo, “**An Efficient Algorithm for Exploiting Multiple Arithmetic Units,**” IBM Journal of R&D, Jan. 1967.
- What is the major difference today?
 - **Precise exceptions:** IBM 360/91 did NOT have this
 - Patt, Hwu, Shebanow, “**HPS, a new microarchitecture: rationale and introduction,**” MICRO 1985.
 - Patt et al., “**Critical issues regarding HPS, a high performance microarchitecture,**” MICRO 1985.
- Variants used in most high-performance processors
 - Initially in Intel Pentium Pro, AMD K5,
 - Alpha 21264, MIPS R10000, IBM POWER5, IBM z196, Oracle UltraSPARC T4, ARM Cortex A15

Two Humps in a Modern Pipeline



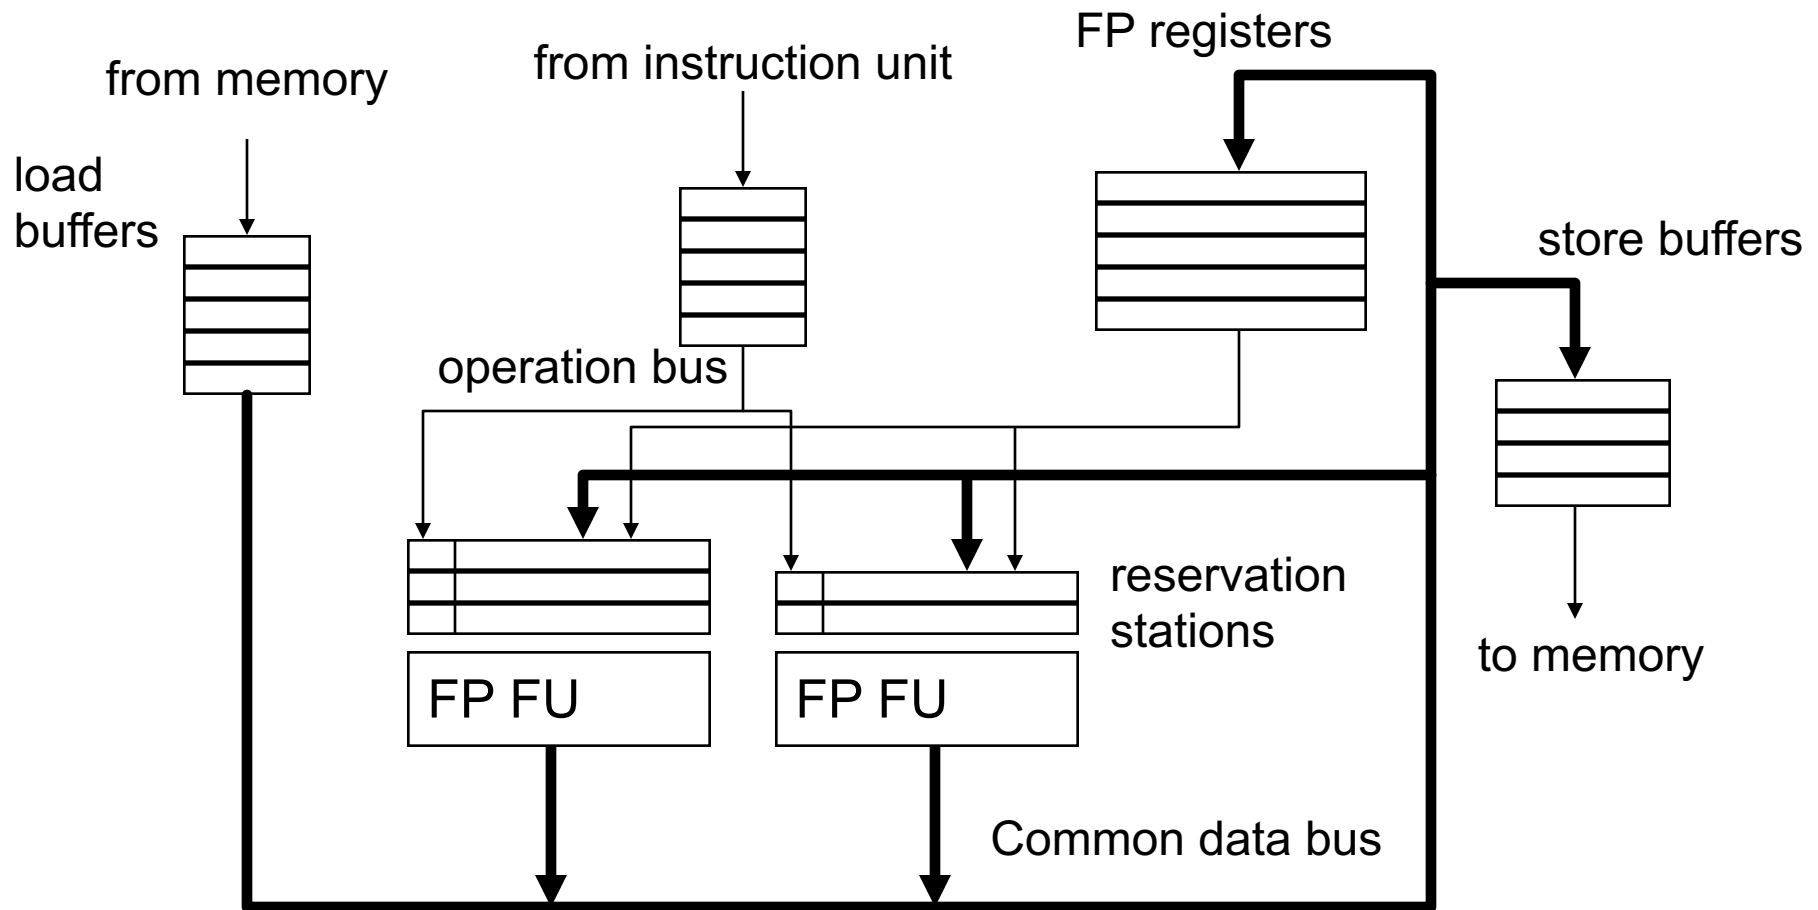
in order

out of order

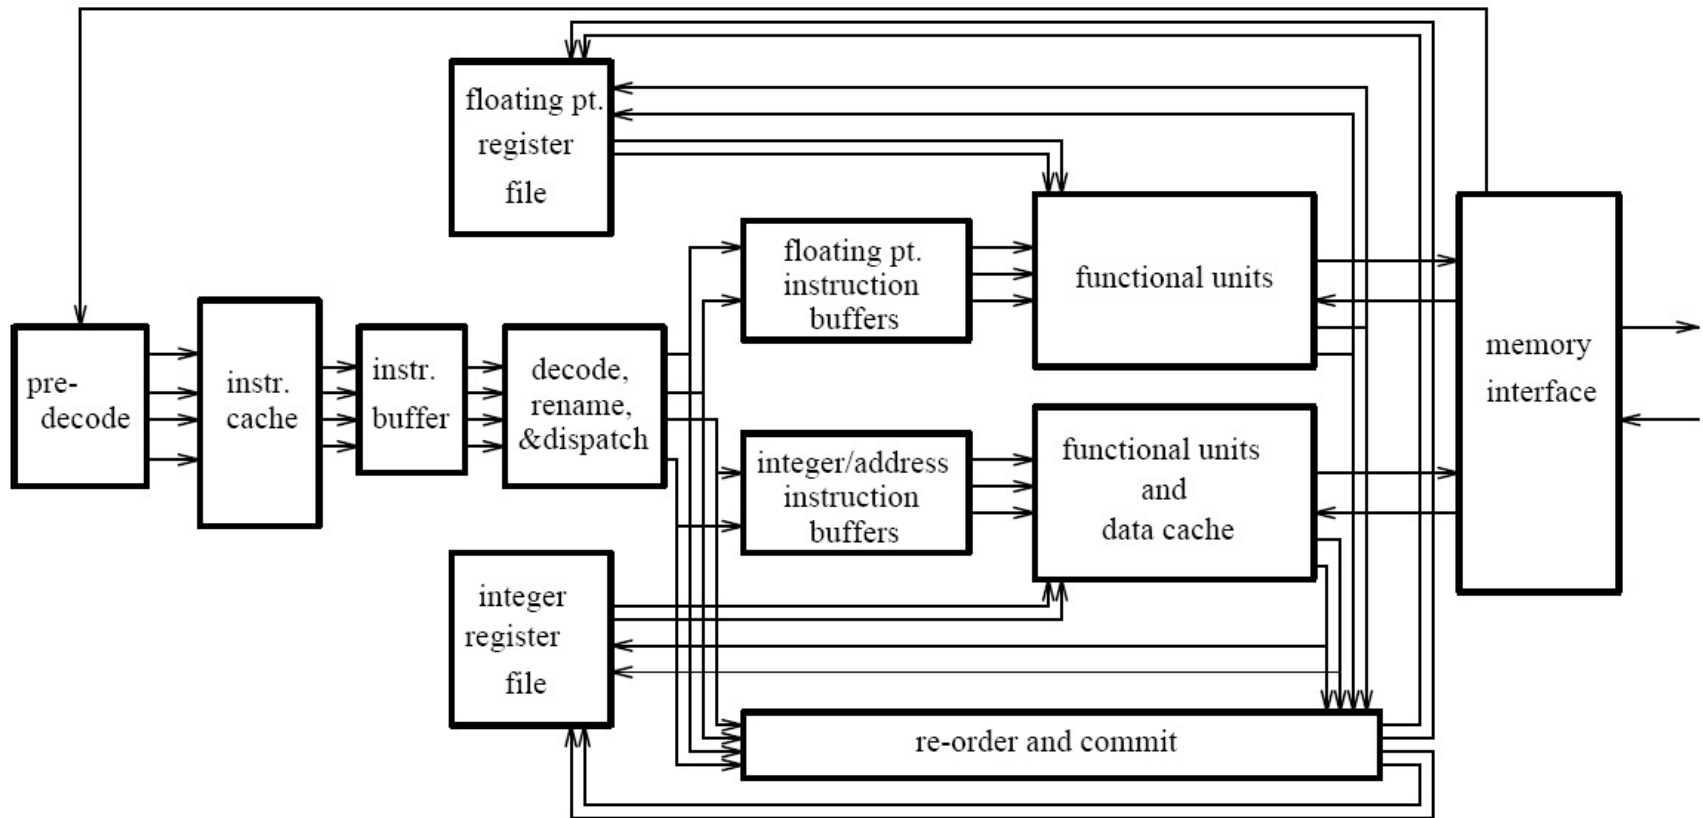
in order

- Hump 1: Reservation stations (scheduling window)
- Hump 2: Reordering (reorder buffer, aka instruction window or active window)

Tomasulo's Machine: IBM 360/91



General Organization of an OOO Processor



Register Renaming

- Output and anti dependencies are not true dependencies
 - **WHY? The same register refers to values that have nothing to do with each other**
 - **They exist because not enough register ID's (i.e. names) in the ISA**
- The register ID is **renamed** to the reservation station entry that will hold the register's value
 - Register ID → RS entry ID
 - Architectural register ID → Physical register ID
 - After renaming, RS entry ID used to refer to the register
- This eliminates anti- and output- dependencies
 - **Approximates the performance effect of a large number of registers even though ISA has a small number**

Tomasulo's Algorithm: Renaming

- Register rename table (register alias table)

	tag	value	valid?
R0			1
R1			1
R2			1
R3			1
R4			1
R5			1
R6			1
R7			1
R8			1
R9			1

Tomasulo's Algorithm

- If reservation station available before renaming
 - Instruction + renamed operands (source value/tag) inserted into the reservation station
 - Only rename if reservation station is available
- Else stall
- While in reservation station, each instruction:
 - Watches common data bus (CDB) for tag of its sources
 - When tag seen, grab value for the source and keep it in the reservation station
 - When both operands available, instruction ready to be dispatched
- Dispatch instruction to the Functional Unit when instruction is ready
- After instruction finishes in the Functional Unit
 - Arbitrate for CDB
 - Put tagged value onto CDB (tag broadcast)
 - Register file is connected to the CDB
 - Register contains a tag indicating the latest writer to the register
 - If the tag in the register file matches the broadcast tag, write broadcast value into register (and set valid bit)
 - Reclaim rename tag
 - no valid copy of tag in system!

An Exercise

MUL R3 \leftarrow R1, R2

ADD R5 \leftarrow R3, R4

ADD R7 \leftarrow R2, R6

ADD R10 \leftarrow R8, R9

MUL R11 \leftarrow R7, R10

ADD R5 \leftarrow R5, R11



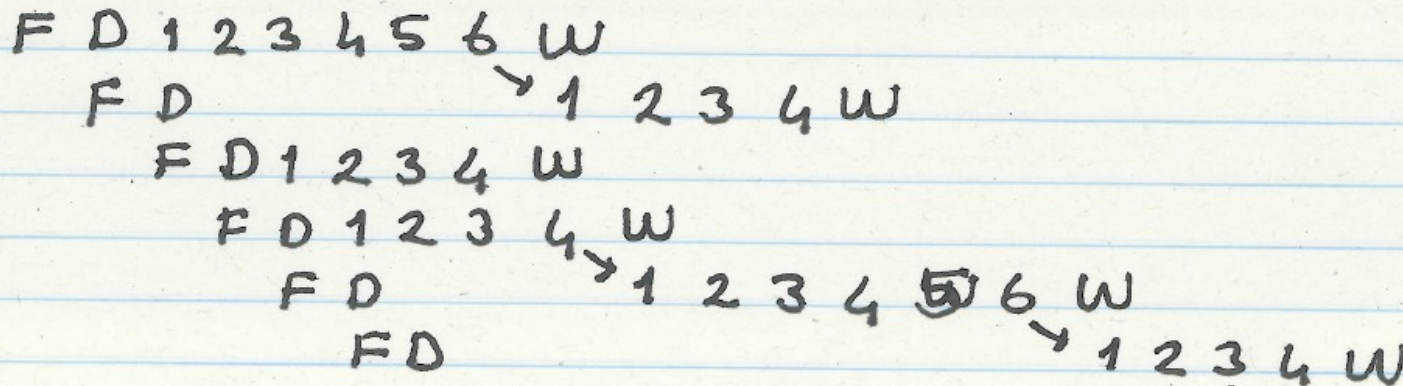
- Assume ADD (4 cycle execute), MUL (6 cycle execute)
- Assume one adder and one multiplier
- How many cycles
 - in a non-pipelined machine
 - in an in-order-dispatch pipelined machine with imprecise exceptions (no forwarding and full forwarding)
 - in an out-of-order dispatch pipelined machine imprecise exceptions (full forwarding)

Stall vs. Forwarding from ROB

OoO execution with ROB

Forwarding

MUL R3 \leftarrow R1, R2
ADD R5 \leftarrow R3, R4
ADD R7 \leftarrow R2, R6
ADD R10 \leftarrow R8, R9
MUL R11 \leftarrow R7, R10
ADD R5 \leftarrow R5, R11



Tomasulo's algorithm + full forwarding

20 cycles

How it works

- RAT (Register Alias Table)

	v?	tag	value
R0	1		
R1	1		
R2	1		
R3	1		



- Reservation Station for Functional Unit (FU)

	v?	tag	value
a	1		
b	1		
c	1		
d	1		

	v?	tag	value
a	1		
b	1		
c	1		
d	1		

	v?	tag	value
x	1		
y	1		
z	1		
w	1		

	v?	tag	value
x	1		
y	1		
z	1		
w	1		



Cycle 0

- RAT all valid
- Let's assume $R0=0$, $R1=1$, $R2=2$, $R3=3$, ...
- Reservation station are all invalid

Cycle2

- Decode MUL R3 \leftarrow R1, R2
- Allocate MUL ALU RS
- R3 is invalid



	v?	tag	value
R1	1		1
R2	1		2
R3	0	x	?
R4	1		

	v?	tag	value
a	1		
b	1		
c	1		
d	1		

	v?	tag	value
a	1		
b	1		
c	1		
d	1		

	v?	tag	value
x	1	~	1
y	1		
z	1		
w	1		

	v?	tag	value
x	1	~	2
y	1		
z	1		
w	1		



Cycle 3

- MUL ALU begins execution R3 invalid
- Decode ADD $R5 \leftarrow R3, R4$



	v?	tag	value
R3	0	x	?
R4	1	~	4
R5	0	a	?
R6	1		

	v?	tag	value
a	0	x	?
b	1		
c	1		
d	1		

	v?	tag	value
a	1	~	4
b	1		
c	1		
d	1		

	v?	tag	value
x	1	~	1
y	1		
z	1		
w	1		

	v?	tag	value
x	1	~	2
y	1		
z	1		
w	1		



Cycle 4

- ADD cannot begin execution (Stall in RS_a)
- MUL consumed 2 cycles
- Decode ADD R7 ← R2, R6



	v?	tag	value
R2	1	~	2
R3	0	x	?
R4	1	~	4
R5	1	a	?
R6	1	~	6
R7	0	b	?

	v?	tag	value
a	0	x	?
b	1	~	2
c	1		
d	1		

	v?	tag	value
a	1	~	4
b	1	~	6
c	1		
d	1		

	v?	tag	value
x	1	~	1
y	1		
z	1		
w	1		

	v?	tag	value
x	1	~	2
y	1		
z	1		
w	1		



Cycle 7

- ADD cannot begin execution (Stall in RS_a)
- MUL consumed 2 cycles
- Decode ADD R7 \leftarrow R2, R6



	v?	tag	value
R2	1	~	2
R3	0	x	?
R4	1	~	4
R5	1	a	?
R6	1	~	6
R7	0	b	?
R8	1	~	8
R9	1	~	9
R10	0	c	?
R11	0	y	?

	v?	tag	value
a	0	x	?
b	1	~	2
c	1	~	8
d	0	a	?

	v?	tag	value
a	1	~	4
b	1	~	6
c	1	~	9
d	0	y	?

	v?	tag	value
x	1	~	1
y	0	b	?
z	1		
w	1		

	v?	tag	value
x	1	~	2
y	0	c	?
z	1		
w	1		



And so on...

- Brief overview of OoO.
 - Fetch in-order, execute OoO, commit in-order
 - Do not wait previous long-running instruction
 - Execute as fast as you can (when no flow-dependency)
 - Register renaming
- Recent trends
 - SIMD, MIMD parallel data flow model
 - Parallel programming many cores utilization
 - For large data handling...
 - 2K cores programming
- Next week, we will move on cache/memory hierarchy